



DESIR: Decoy-Enhanced Seamless IP Randomization

Jianhua Sun, Kun Sun

Department of Computer Science
College of William and Mary



Roadmap

- Introduction
- Threat Model
- System Architecture
- Implementation
- Security Analysis
- Performance Evaluation
- Conclusion



Roadmap

- Introduction
- Threat Model
- System Architecture
- Implementation
- Security Analysis
- Performance Evaluation
- Conclusion



Background

- Network reconnaissance attacks have been effective due to the **static** nature of current network and system configurations
- Existing IP randomization based solutions shift network attack surface, including:
 - IP and MAC addresses, open ports, network topology



Limitations of previous approaches

- Effectiveness of IP randomization is reduced due to the small number of alive IP addresses at one time
 - Small security entropy
- Existing active connections may be disrupted when the IP addresses of the servers are changed
 - Negative impact on user experience



Contribution Highlights

Solve two major challenges!

- **Service Security** against malicious users
 - Fortify IP randomization with a large number of decoys that shuffle their address along with the real servers
- **Service Availability** to legitimate users
 - Develop a seamless network connection migration mechanism to keep alive the pre-existing connections



Roadmap

- Introduction
- Threat Model
- System Architecture
- Implementation
- Security Analysis
- Performance Evaluation
- Conclusion



Threat Model and Assumptions

- Focus on **persistent reconnaissance** attacks
 - Not consider insiders that deliberately disclose the current server IP address to attackers
- Adversaries are not in the same subnet with legitimate users
 - Cannot obtain the server IP addresses through packet eavesdropping
- Secret keys are shared between the legitimate users and the servers
- The protected network consists of a large number of IP addresses to accommodate decoy nodes




Roadmap

- Introduction
- Threat Model
- System Architecture
- Implementation
- Security Analysis
- Performance Evaluation
- Conclusion



System Architecture





Randomization Controller

- Decision engine
 - Determine randomization frequency, choose algorithm to generate new network configurations
- Configuration generator
 - Control overall topology and regenerate new configuration
 - Guarantee there is no interference in IP address assignment
- Migration console
 - Distribute the new configuration to the servers and decoy subsystem



Decoy Bed

- Communication module
 - Receive new configuration settings from the randomization controller
 - Determines the overall architecture of the decoy network
- Decoy generator
 - Regenerates the decoy network
 - Flexible to deploy both high-interaction and low-interaction decoys



Migration Module

- Connection interception
 - Introduce a pair of internal and external addresses to detach transport layer identity from network layer identity
- Connection translation
 - Intercepts packets in the network layer and translates the internal addresses in the packet headers to or from the external addresses for outgoing/incoming packets
- Connection migration
 - Coordinates the moving of server associated with active connections to another IP address




Roadmap

- Introduction
- Threat Model
- System Architecture
- Implementation
- Security Analysis
- Performance Evaluation
- Conclusion



VM-based implementation

- **Host configuration**
 - Ubuntu 14.04 and KVM
 - Intel Core i7-4712HQ CPU, 16GB RAM
- **Five VMs:** decoy bed, real server, AS, client, attacker
 - Each VM allocated one host CPU and 2GB memory
 - Decoy VM runs Ubuntu 12.04
 - Remaining VMs run Fedora 15 with Linux kernel 2.6.38





Three-level decoy bed

- Virtual machine level
 - Virtual machines with fully functional OS and applications
- Operating system level
 - Containers deployed using LXC in Honeybed VM
- Process level
 - Honeyd deployed in containers



Seamless connection migration

- Connection interception
 - Intercept system calls for connection setup from the application layer to transport layer
 - For TCP connection, socket, accept, connect, close, getsockname, getpeername
 - For UDP connection, send/recv
- Connection translation - Use `iptables` to do NAT
 - Client side: DNAT on OUTPUT chain, SNAT on POSTROUTING chain
 - Server side: DNAT on PREROUTING chain, SNAT on INPUT chain
 - Use `mangle` table to block connection attempt to internal addresses



Seamless connection migration

- Connection migration
 - Use two daemons within both endpoints to negotiate with each other the migration based on a predefined protocol
 - i. Suspend the connection
 - ii. Restore after IP randomization is finished; create a VIF to which the internal connection is attached
 - Synchronously randomize the communication ports
 - Encrypt the negotiation messages with a shared secret key



Roadmap

- Introduction
- Threat Model
- System Architecture
- Implementation
- Security Analysis
- Performance Evaluation
- Conclusion



Scanning unaware of IP randomization defense

- Static IP address layout, no need to scan a single IP twice
 - **Sampling without replacement** problem
 - Expected number of probes to identify real server in a n IP pool:
$$(n+1)/2$$
- Randomize the IP space after each probe. If the attacker takes a single-round scan, the expected number of probes is $(1 - 1/e) n = 0.63n$
- Attacker needs to pay 26% more efforts to locate the real server



Scanning aware of IP randomization defense

- Identifying the target server IP can be treated as a **sampling with replacement** problem
 - No matter whether the IP space is periodically re-randomized or not
- The number of probes m performed is a geometric random variable with probability $p=1/n$
- Therefore, the expected number of probes is $1/p=n$




Roadmap

- Introduction
- Threat Model
- System Architecture
- Implementation
- Security Analysis
- Performance Evaluation
- Conclusion



Microbenchmark - delay overhead

- Use sftp, ftps, tftp to transfer 1GB file
- sftp/ftps: 1% - 9% overhead
- tftp: 5% - 15% overhead







System Overhead

- Use Netperf to measure migration related overhead in terms of network latency and throughput
- Three configurations
 - **Vanilla** - a stock Linux with Netfilter firewall rules loaded on boot
 - **Vanilla+Virt** - system with both Netfilter and migration module loaded, where the connections are not migrated but the socket system calls are intercepted
 - **Migration** - with both Netfilter and migration module loaded and all connections are migrated



System overhead

- **Vanilla+Virt** - connection interception incurs no overhead
- **Migration** - 2% to 7% overhead - incurred by connection translation





System Scalability

Connection migration overhead breakdown

Total number of connections	Total suspension time (s)	Suspension time per connection (ms)	Total restoration time (s)	Restoration time per connection(ms)	Memory consumption (KB)	
					Virtual interface	Address mapping
10	0.14	13.77	0.35	35.31	10.63	0.86
50	0.69	13.81	1.83	36.65	53.13	4.30
100	1.45	14.51	3.74	37.43	106.25	8.59
500	7.33	14.67	17.37	34.74	531.25	42.97
5000	73.5	14.7	174.1	34.82	5315	429.8

- Average time to suspend a connection: 14 ms; to restore: 35 ms
- Virtual interface accounts for 90% of memory consumption, 1.06 KB each
- 5 MB memory overhead when migrating 5000 connections



Roadmap

- Introduction
- Threat Model
- System Architecture
- Implementation
- Security Analysis
- Performance Evaluation
- Conclusion



Conclusion

- We propose a defense framework for constructing a dynamically mutable network with a number of decoys to protect the real servers against scanning attacks
- Our solution can ensure seamless connection migration with IP address randomization and guarantee both service availability and service security of the real servers
- We implement a VM-based prototype, which shows that our system has good scalability and acceptable network and system performance overhead